

Variable quality image compression system based on SPIHT

A. Järvi^{a,b,1} J. Lehtinen^{a,b,2} O. Nevalainen^b

^a *Turku Centre for Computer Science, Lemminkäisenkatu 14 A, FIN-20520 Turku,
Finland*

^b *Department of Computer Science, University of Turku, Lemminkäisenkatu 14 A,
FIN-20520 Turku, Finland*

¹ Corresponding author. E-mail ajarvi@cs.utu.fi, voice +358 2 3338795, fax +3582
2410154

² Joonas Lehtinen acknowledges support by the Academy of Finland.

Number of pages is 30.

Number of figures is 10.

Number of tables is 1.

Keywords: Image compression, SPIHT, Wavelet transform, Region of interest, variable quality image compression, VQIC.

Abstract

An algorithm for variable quality image compression is given. The idea is to encode different parts of an image with different bit rates depending on their importance. Variable quality image compression (VQIC) can be applied when *a priori* knowledge on some regions or details being more important than others is available. Our target application is digital mammography, where high compression rates achieved with lossy compression are necessary due to the vast image sizes, while relatively small regions containing signs of cancer must remain practically unchanged. We show how VQIC can be implemented on top of SPIHT [11], an embedded wavelet encoding scheme. We have revised the algorithm to use matrixes, which gives more efficient implementation both in terms of memory usage and execution time. The effect of the VQIC on the quality of compressed images is demonstrated with two test pictures: a drawing and a more relevant mammogram image.

1 Introduction

The number of large digital image archives is increasing rapidly in many fields, including health care. The cost efficient archiving causes a need for high quality image compression techniques aiming at major savings in storage space and network bandwidth when transmitting the images. Despite the potentially critical nature of medical images, some degeneration of the image quality must be allowed, since the best *lossless compression* methods can only about half the size of a typical medical image. For example a digital mammogram with pixel size of $50\mu\text{m}$ is approximately of size 5000x5000 pixels with 12 bits per pixel, and thus needs about 50 Mb for storage without compression. This clearly demonstrates the need for *lossy compression*.

Lossy image compression methods are usually designed to preserve perceived image quality by removing subtle details, that are difficult to see with human eye. The frequent quality measure used for evaluation of the distortion in a compressed image is the *mean square error* (MSE). However, in medical imaging, the distortion of an image is defined as the impact that compression causes to diagnostic accuracy, and finally to clinical actions taken on the basis of the image [1]. In this application area, there is an evident conflict between the two opposite goals; achieving high compression ratio and maintaining diagnostically lossless reconstruction accuracy. One possible way to alleviate this conflict is to *design an image compression method that uses lossy compression, but saves more details in important regions of the image than in other regions*. General-purpose image compression methods can also be considered to fit into this scheme; important regions and details are those, that the human visual

system is sensitive to. In medical imaging, the definition of important regions involves application specific medical knowledge. Obviously the accuracy of this knowledge is crucial for good performance of the system. If the criteria for important regions are too loose, the gain in compression ratio is lost. On the other hand, with too strict criteria the quality requirements are not met, since some important regions are treated erroneously as unimportant.

Today, the standard in lossy image compression is the JPEG [2] algorithm, which is based on the scalar quantization of the coefficients of the windowed Discrete Cosine Transforms (DCT), followed by entropy coding of the quantization results. JPEG is generally accepted and works well in most cases, but because it uses DCT and divides the image into blocks of fixed size, it may distort or even eliminate small and subtle details. This can be a serious drawback in digital mammography, where images contain a large number of diagnostically important small low contrast details, that must preserve their shape and intensity.

Wavelet based image compression methods are popular and some of them can be considered to be “state of the art” in general-purpose lossy image compression (see [6] for an introduction to the topic). Wavelet compression methods can be divided into three stages: *wavelet transform*, *lossy quantization* and *encoding* of the wavelet coefficients and *lossless entropy coding* [14]. The wavelet transform is used to de-correlate the coefficients representing the image. The transform collects the image energy to relatively small number of coefficients, compared to the original highly correlated pixel representation. In the quantization phase this sparse representation and dependencies between coefficients are exploited with specially tailored quantization and coding schemes. The widely known *embedded zerotree encoding* (EZW) by Shapiro [12] is an excellent example of such coding scheme, and also a good reference to wavelet based image compression in general.

One of the most advanced wavelet based image compression techniques is SPIHT (Set Partitioning In Hierarchical Trees) by Said and Pearlman [11]. SPIHT is clearly a descendant of EZW using similar zerotree structure and *bitplane coding*. In bitplane coding bits of wavelet coefficients are transmitted in the order of their importance, i.e. the coefficients are encoded gradually with increasing accuracy. Because of this, the encoding can be stopped at any stage. The decoder then approximates the values of the original coefficients with precision depending on the number of bits coded for each coefficient. This property called *embedded coding* is the main reason for choosing SPIHT as the basis of the *Variable Quality Image Compression system* (VQIC). The implementation of variable quality property is straightforward in embedded coding: the encoding of the coefficients of the whole image is ceased somewhere in the middle, and subsequently only bits of coefficients that influence the important regions are encoded.

Another reason for choosing SPIHT is that it appears to perform very well in terms of compression ratio; in the case of digital chest X-rays, a compression ratio of 40:1 has been reported to cause no significant difference when compared to the original in a radiologist evaluation [4]. In an other study SPIHT compression of mammograms to ratio 80:1 has been found to yield image quality with no statistically significant differences from the original mammogram [9]. A further indication of good performance is the fact that the output of SPIHT encoding is so dense, that additional compression with lossless entropy coding gives an extra gain of only few percentages [11].

VQIC technique can be applied to any image that is spatially segmentable to a set of regions that must be saved in better quality. We do not cover the segmentation problem in this paper since it is completely application specific. Since we are targeting at applications where important regions are small, the segmentation is done with some kind of feature detector. The feature detection task for VQIC purpose is considerably easier compared to applications, where the interest is in the presence or absence of the feature. In VQIC, a moderate number of false positive detections can be tolerated, as long as all of the true features are detected. Thus most existing feature detection algorithms are suitable, because they can be tuned to be oversensitive. Several suitable fully automatic segmentation methods exist for this purpose in the field of medical imaging [5], especially for mammography [3].

Before describing further details, we briefly discuss relevant research. VQIC has been used in the compression of image sequences in video conferencing [10]. In this work, the pixel values are predicted and the prediction errors are transformed by 2D-DCT. Coefficients of the blocks with minor importance are quantized with coarser level than more important details, heads and shoulders. In another research, the importance of a region is determined on the basis of the visibility of distortions to human eye [8]. This information is used in the construction of a constant quality MPEG stream by adjusting quantization parameters defined by the MPEG standard. The focus in both papers is the segmentation of important regions, and VQIC is achieved by variable quantification of DCT coefficients. A recent paper by D. Shin, H. Wu and J. Liu describes a selective compression technique, which integrates detection and compression algorithms into one system [13]. The compression technique used is intraband coding of wavelet packet transform coefficients, where variable quality is achieved by scaling the wavelet coefficients in important regions to increase their priority in coding. The method is tested with a digital mammogram, where detected microcalcifications are considered as ROIs. Even though the aims of this work are close to ours, the actual methods differ considerably.

Our work is organized as follows. In section two we introduce an algorithm called *variable quality SPIHT* (*vqSPIHT*), which is basically a reimplementation of SPIHT, with the added VQIC functionality. We revise the memory

organization of SPIHT to use matrix-based data structures. This new implementation reduces the working storage requirements of the algorithm considerably. In section three we discuss the compression performance of vqSPIHT algorithm, and show with two examples that it can be superior to SPIHT or JPEG. We first demonstrate this with a set of details in a high contrast drawing compressed to several bit rates with all three compression algorithms. We also discuss a more relevant application for VQIC – compression of digital mammograms, and make a comparison between SPIHT and vqSPIHT compressed images.

2 vqSPIHT algorithm

In this section we explain informally the basic ideas behind SPIHT and vqSPIHT algorithms to facilitate the reading of the vqSPIHT algorithm in a pseudo-code format. We also discuss the implementation based on matrix data structures and its implications to practical memory requirements.

2.1 Structure of the wavelet transformed coefficient table

Wavelet transform converts an image into a coefficient table with approximately the same dimensions as the original image. Fig. 1 shows the structure of the *wavelet coefficient table*, which contains three wavelet coefficient pyramids (pyramids *A*, *B* and *C*) and one table of *scaling coefficients S*. The scaling coefficients represent roughly the mean values of larger parts of the image and wavelet coefficient details of various sizes. Since in practice the transform is stopped before scaling table *S* would shrink to a single coefficient, table *S* looks like a miniature version of the original image. The top levels of the three wavelet pyramids are located adjacent to the scaling table (level three in Fig. 1), and contain coefficients representing large details, whereas coefficients at level zero contribute mainly to the smallest details in the image. Pyramid *A* contains coefficients for vertical details and pyramid *B* respectively for horizontal details. The third pyramid *C* contains correction coefficients needed in the reconstruction of the image from pyramids *A*, *B* and table *S*.

In the SPIHT algorithm, the pyramids are divided into *sub-pyramids* which corresponds to zerotrees in EZW. A sub-pyramid has a top element somewhere in the coefficient table, and contains four coefficients one level lower in the corresponding spatial location in the same pyramid, 16 elements two levels lower, and so on. The sub-pyramids are extended to the scaling coefficients *S* in the following way. The scaling coefficients are grouped into groups of four.

The coefficient in the upper left corner has no descendants, whereas the three remaining coefficients in the group (upper right corner, lower left corner and lower right corner) serve as top elements of three sub-pyramids in pyramids A , B and C in corresponding order.

In our approach to VQIC we must determine which coefficients contribute to the value of a given pixel in the original image. In an *octave-band decomposition*, which we use, a coefficient of any pyramid in higher level corresponds to four coefficients in the next level at same spatial location. We use this rule of multiples of four in choosing important coefficients.

2.2 The basic functioning of SPIHT

To understand how SPIHT works, one must keep in mind that it is not a complete image compression scheme, but a method tailored for optimal embedded encoding of wavelet transformed image coefficients. The encoding is optimal in the sense of MSE. SPIHT does not presuppose any particular wavelet transform. The only requirement is that the transform has the octave band decomposition structure, as described above. Also, the optimal encoding with respect to MSE is achieved only if the transform is unitary, which is the case in *orthogonal and biorthogonal* wavelet transforms [14]. In the implementation of vqSPIHT, we use biorthogonal B97 wavelets [14].

2.2.1 Bitplane coding in SPIHT

Optimal progressive coding of SPIHT is implemented with a *bitplane coding scheme*. The order of coding is based on the energy saving property of unitary transforms, here the wavelet transform. This property states that the larger the wavelet coefficient is, the more its transmission reduces the MSE. Furthermore, since SPIHT uses uniform scalar quantization, transmission of a more significant bit in any coefficient reduces the MSE more than transmission of a less significant bit in a possibly larger coefficient [11].

According to this principle, all coefficients are sorted to a decreasing order by the *number of significant bits*. The number of significant bits in the coefficient having the largest absolute value is noted by n . The output is generated by transmitting first all the n :th bits in coefficients that have at least n significant bits, then $(n - 1)$:th bits of coefficients that have at least $(n - 1)$ significant bits, and so on. Because the most significant bit of a coefficient is always one, the sign of the coefficient is transmitted in place of the most significant bit.

In addition to transmitted bitplanes, the sorting order and the length of each bitplane are needed in the decoder to resolve the location of each transmitted

bit in the reconstructed wavelet coefficient table. This information is not transmitted explicitly, instead the same algorithm is used in both the encoder and the decoder, and all branching decisions made in the encoder are transmitted to the decoder. The branching decisions are transmitted interleaved with the bitplane coded bits and the signs of coefficients. Because of progressive nature of SPIHT coding, the transmitted bit stream can be truncated at any point and the original coefficient matrix approximated with optimal accuracy with respect to the number of transmitted bits. For each coefficient, the most significant not transmitted bit is set to one, and the rest to zero, thus achieving a good approximation in uniform quantization.

2.2.2 Exploitation of the pyramid structure of wavelet transform

An important property in most natural images is that the low- and high frequency components are spatially clustered together. This means in the wavelet coefficient pyramid, that there is high correlation between the magnitudes of the coefficients of different levels in corresponding locations. Also, since the variance of the frequency components tends to decrease with increasing frequency, it is very probable that the coefficients representing fine details in a particular spatial location will be small, if there is a region of small coefficients in the corresponding location on coarser level of the pyramid. Thus it is probable that there exists sub-pyramids containing only zeroes on the current bitplane. These *zerotrees* can be encoded with one bit, thus cutting down the number of sorting decisions considerably and also the branching decisions that must be transmitted to the decoder. The way these dependencies between coefficients are exploited in SPIHT coding is described in the presentation of vqSPIHT algorithm.

2.3 The vqSPIHT algorithm

2.3.1 Extension of SPIHT to vqSPIHT

To expand the SPIHT algorithm to vqSPIHT, we define a *Region Of Interest* (ROI) as a region in the image that should be preserved in better quality than the rest of the image. ROIs can be presented as a binary map that is highly compressible with simple run-length encoding and thus does not affect bit rate significantly.

In selective coding mode of vqSPIHT, only coefficients affecting ROIs are coded. This mode is triggered when a certain *percentage α of the wanted final output file size* has been reached. The choice of α is important, and its best value is highly application dependent. Some applications might demand more

sophisticated definition of α depending on the file size, area of the ROIs and some indicator on how ROIs are scattered, for example.

To implement selective coding, we construct a *look-up-table* (LUT), that is used in the function $influence(i, j)$ defining whether a coefficient (i, j) contributes to any ROI or not. The LUT is constructed by scaling each level of all three pyramids to the same size as the original image, and comparing the map of ROIs to the scaled levels. All the coefficients that overlap with any ROI are marked in the LUT.

2.3.2 Implementation with matrices

Instead of lists that are used in the original implementation of SPIHT, our implementation of vqSPIHT uses two matrices for keeping track of significant and insignificant coefficients and sub-pyramids. With the matrix data structures we can considerably reduce the working storage requirements of encoding and decoding.

We introduce a *Point Significance Matrix* (PSM) to indicate whether a coefficient is known to be *significant*, *insignificant* or still has an *unknown* state. The labels of PSM are coded with two bits and the dimensions of the PSM are the same as in the coefficient table. We also need a *Sub-Pyramid List Matrix* (SPLM), which is used for maintaining an implicit list of the sub-pyramids containing only insignificant coefficients. The list structure is needed, because the order of sub-pyramids must be preserved in the sorting algorithm. The dimensions of the SPLM are half of the dimensions of the coefficients matrix, because the coefficients on the lowest level of the pyramids can not be top elements of sub-pyramids. There are two types of sub-pyramids. A sub-pyramid of type *A* contains all descendants of a particular coefficient, excluding the top coefficient itself. A sub-pyramid of type *B* is otherwise similar, but the immediate offspring of the top coefficient is excluded in addition to the top coefficient. The list structure with SPLM is simple: the lower bits of an element tells the index of the next element in the list. The type of the sub-pyramid is coded with the highest bit.

2.3.3 Pseudo-code of vqSPIHT

The algorithms for the encoder and the decoder are similar. We use notation “input/output x ”, which consists of two steps: In the encoder x is first calculated and then transmitted to entropy coder; in the decoder x received from decoder and then used in the construction of a new estimate for the coefficient. Variable nbc indicates the number of bits transmitted or received thus far. Constant $filesize$ indicates the requested final file size in bits.

Function $influence(i, j)$ is defined to be *true*, if the element (i, j) in the LUT is marked to influence an ROI, and *false* otherwise. Let $c_{i,j}$ be the value of the coefficient (i, j) in the wavelet coefficient table and the coordinate pair (i, j) denote either a single coefficient or a whole sub-pyramid of type *A* or *B* having coefficient at (i, j) as the top element. The meaning of (i, j) will be evident from the context. Finally we define the significance of a coefficient with function $S_n(c_{i,j})$ as follows: $S_n(c_{i,j}) = 1$ if the number of bits after the first 1-bit in the absolute value of $c_{i,j}$ is at least $n - 1$, otherwise $S_n(c_{i,j}) = 0$. The significance of a sub-pyramid is defined with function $S_n(i, j)$. $S_n(i, j) = 0$ if $S_n(c) = 0$ for all coefficients c belonging to sub-pyramid (i, j) , otherwise $S_n(i, j) = 1$.

The input for both the encoder and decoder is the wavelet coefficient table, the ROI map, α and *filesize*. The output of the decoder is an approximation of the original wavelet coefficient table. See Fig. 2 for an outline of the algorithm.

1. Initialization

- Input/output n , which is the number of significant bits in the coefficient having the largest absolute value.
- Construct LUT according to the given ROIs.
- Set $nbc = 0$.
- Set the PSM label of all scaling coefficients (coefficients in the area S in Fig. 1) to *insignificant* and the PSM label of all other coefficients to *unknown*.
- Create a list of all scaling coefficients that have descendants in SPLM and make them of type *A*.

2. Sorting step for PSM

- For every element (i, j) in PSM do
 - ★ If $(nbc/filesize < \alpha \text{ OR } influence(i, j))$ then
 - If (i, j) is labeled to be *insignificant* do:
 - Input/output $S_n(c_{i,j})$.
 - If $S_n(c_{i,j}) = 1$ then set the PSM label (i, j) to *significant* and input/output the sign of $c_{i,j}$.
 - Else If (i, j) is labeled to be *significant* do:
 - Input/output the n -th most significant bit of $|c_{i,j}|$.

3. Sorting step for SPLM

- For each element (i, j) in the list in SPLM do:
 - ★ If sub-pyramid (i, j) is of type *A* AND $(nbc/filesize < \alpha \text{ OR } influence(i, j))$ then
 - Input/output $S_n(i, j)$.
 - If $S_n(i, j) = 1$ then

- + For each (k, l) belonging to immediate offspring of (i, j) do:
 - * If $(nbc/filesize < \alpha \text{ OR } influence(k, l))$ then
 - Input/output $S_n(c_{k,l})$.
 - If $S_n(c_{k,l}) = 1$ then set the PSM label (k, l) to *significant* and input/output the sign of $c_{k,l}$, else set the PSM label (k, l) to *insignificant*.
- + If (i, j) is not on one of the two lowest levels of the pyramid then move (i, j) to the end of the list in SPLM and change its type to *B*, else remove (i, j) from the list in SPLM.
- * If sub-pyramid (i, j) is of type *B* AND $(nbc/filesize < \alpha \text{ OR } influence(i, j))$ then
 - Input/output $S_n(i, j)$.
 - If $S_n(i, j) = 1$ then
 - + For each (k, l) belonging to immediate offspring of (i, j) do:
 - If $(nbc/filesize < \alpha \text{ OR } influence(k, l))$ then add (k, l) to the end of list in SPLM as sub-pyramid of type *A*.
 - + Remove (i, j) from the list in SPLM.

4. Quantization-step update

- o If $n > 0$ then
 - Decrement n by 1.
 - Jump to the beginning of the PSM sorting step 2.

The *original SPIHT* algorithm always transmits bits in two alternating phases: In the first phase the branching decisions of the sorting step and the signs of new significant coefficients are transmitted. In the second phase the bits of all significant coefficients on the current bitplane are transmitted. Interrupting the first phase can cause transmission of branching decisions that can not be used in reconstruction. In *vqSPIHT*, we have therefore combined the sending of the significant coefficients to the sorting phase to avoid the problem.

3 Test results

3.1 Numerical quality indicators

As a measure of image quality, we use the *Point Signal to Noise Ratio* (PSNR) for the whole image and for the ROIs:

$$D_{PSNR} = 10 \log_{10} \frac{2^{\text{bpp}} - 1}{D_{MSE}} dB.$$

It is well-known that PSNR does not give objective estimate of image quality,

but it correlates with the amount of distortion in the image and it can be used for comparing the quality of images compressed with algorithms causing similar distortion. As a measure for the amount of compression we use the number of *bits per pixel* (bpp).

3.2 The comic test image

The vqSPIHT algorithm was constructed for the needs of digital mammography. However, mammograms are rather smooth and thus easily hide compression artifacts. To better illustrate the effect of VQIC, we use a comic picture of size 420x480 with 8 bpp as the first test image, Fig. 3. GIF¹, JPEG, SPIHT and vqSPIHT algorithms are used to compress the image having three ROIs covering 2.4% of the image marked on the Fig. 3. The compression results are presented in Table 1 and Fig. 4.

As seen in Table 1, the PSNR of SPIHT and vqSPIHT are similar. With less compression (large bpp), JPEG is also comparable in terms of PSNR, but its visual quality decreases rapidly with decreasing values of bpp (Fig. 4). In JPEG, the resulting file size can not be specified exactly in advance, and thus the bpp values of JPEG are slightly different from those of SPIHT and vqSPIHT. There are no big differences in the performance of these techniques, when only the overall PSNR is evaluated.

When considering the PSNR of ROIs, the situation changes radically. SPIHT performs significantly better than JPEG, but the improvement achieved with vqSPIHT is even greater. We have used quite high α values: 80% and 90% of the size of the output file. Even with these values, PSNR in the ROIs is considerably lower in the vqSPIHT compressed images than in the SPIHT compressed images, while good overall quality (PSNR of whole image) is still maintained. This is partly due to the fact that the coefficients that influence the ROIs also contribute to the areas outside the ROIs. Thus the overall image quality is still improving outside the ROIs after the trigger value α has been reached.

Fig. 4 shows a 88x66 pixel region taken from the comic image and compressed with JPEG, SPIHT and vqSPIHT with different bpp values. The original part of the image is shown in the lower right corner. The selected part includes an ROI, marked on the original image.

The first row of Fig. 4 shows the limit bpp value, where JPEG clearly fails to produce acceptable quality. The image compressed to 0.50 bpp is still recognizable, but the 0.25 bpp image is not. Even the 1.00 bpp image compressed with JPEG has high-frequency noise around the sharp edges. In the 0.5 bpp and 0.25 bpp images the blocking effect introduces additional artifacts. In the

1.00 bpp SPIHT image there is no high-frequency noise. When the bpp value gets smaller, the image gets smoother, and it thus loses small high-frequency details. However, even the 0.25 bpp SPIHT image is recognizable.

The overall image quality of the 90% and 80% 1.00 bpp vqSPIHT images is very close to that of the 1bpp SPIHT image. The visual quality of the 0.50 bpp SPIHT image is similar to the 0.25 bpp vqSPIHT ($\alpha = 80\%$) image on the ROI. Note that the ROI of the 0.10 bpp vqSPIHT image is visually better than the ROI on the 0.25 bpp JPEG image, and of comparable quality with ROI on the 0.25 bpp SPIHT image. In this image, 0.25 bpp corresponds to compression ratio 32:1. It should be noted that SPIHT is designed to perform well on natural images. A comic drawing is a difficult case for SPIHT and thus also for vqSPIHT.

The performance of the vqSPIHT was good when ROI covered only 2.4% of the picture. With the increase of ROI α must decrease to compensate the larger number of coefficients in the ROI in order to maintain the same quality. Because bits are coded in the order of their importance, the bits used in coding of ROI can add considerably less to whole image PSNR than the bits outside ROI. As seen in Fig. 5, the benefits of VQIC rapidly disappear with large ROI.

3.3 The mammogram test image

The second test image, Fig. 6, is a mammogram of size 2185x2925 with 12 bpp. The mammogram test image has been compressed only with vqSPIHT. However the setting of α to 100% makes vqSPIHT function similarly to SPIHT.

In this example, we assume that the micro-calcifications are the only important diagnostic details of a mammogram, that are easily lost in compression. Note that in a study of the applicability of vqSPIHT to digital mammograms also other signs of cancer, like stellate lesions and nodules, should be considered. A micro-calcification location map, shown in Fig. 7, was generated with a micro-calcification detection algorithm slightly modified from the morphological segmentation algorithm of Dengler, Behrens and Desaga [3]. The detection was tuned to be over-sensitive to make sure that all micro-calcifications were detected. Because of this, the algorithm detected also a large number of false calcifications, including the skin-line of the breast. In this test case, there were 323 ROIs covering five percent of the whole mammogram.

We used the bpp values 0.05, 0.10, 0.15, 0.25, 0.50, 0.75, 1.00 and let α take values of 30, 40, 50, 60, 70, 80, 90 and 100 percent of the resulting file size. Fig. 8 shows the PSNR of the whole image as a function of α and the bpp. Lowering α decreases the PSNR of the whole image, but the effect remains moderate with reasonable α values.

Fig. 9 shows the PSNR calculated only on the ROI as a function of α and the bpp. The benefits of the VQIC on ROIs is clearly seen in comparison with the Fig. 8. To point out, the PSNR of ROIs in 1.00 bpp mammogram jumps from 40.29dB to 54.87dB when α decreases from 100% (i.e. SPIHT) to 80%. This causes a very moderate change in the PSNR of the whole image, which decreases from 38.98 to 38.38.

Fig. 10 shows a region containing a micro-calcification cluster taken from a mammogram, that has been compressed using various bpp and α values. A visual comparison between SPIHT and vqSPIHT shows that a mammogram can be compressed to a significantly lower bpp value with vqSPIHT than with SPIHT ($\alpha = 100\%$) to achieve similar preservation of micro-calcifications in the ROIs. The region in the upper left corner has been compressed with SPIHT to compression ratio 12:1. Even with this rather modest compression, a comparison with the original (lower left corner) reveals that the edges of the calcifications have become blurred, some small calcifications have disappeared and some have merged together. When keeping the same bpp 1.00, we notice that setting $\alpha = 70\%$, the micro-calcifications are virtually indistinguishable from the original. With this choice of α , the PSNR of whole image decreases from 38.98dB to 37.51dB. Now, keeping $\alpha = 70\%$, the bpp value 0.15 (compression ratio 1:80) gives a visually comparable reconstruction to the 1.00 bpp SPIHT image (compression ratio 1:12). In this case, the PSNR of the whole mammogram decreases to 34.08dB. This is, however, virtually same as the PSNR of SPIHT 0.15 bpp compressed image, which is 34.10dB.

3.4 Practical memory requirements of the implementation

We first implemented the algorithm using the list data structures of the original SPIHT algorithm [11], but found that this required a large amount of internal memory. The amount of memory needed was very dependent on the values of bpp and α . Typically, the compression of a 12 MB mammogram required at least 120 MB of internal memory during encoding, but with some combinations of bpp and α the memory requirement was considerably larger. The memory is mainly used for representing the coefficient table and the lists that are constantly scanned through. Thus paging the memory to hard disk increases the execution time drastically. Memory requirements can be made independent of the bpp-ratio and α by reimplementing the algorithm using the matrix data structures presented previously. The working memory space dropped to about 50MB and about 40% of that could be paged to disk without significant increase of the execution time. All of the needed memory could be allocated once, which made the memory management efficient in comparison to the slow per-node dynamical memory management of explicit list structures.

4 Summary and conclusions

The idea of VQIC is to use more bits for important details at the cost of unimportant details such as noise. The compression method can be applied in applications where certain small regions in the image are especially important. We have shown that in our target application, compression of digital mammograms, the variable quality compression scheme can improve the compression efficiency considerably. The variable quality property has been integrated into SPIHT, which is one of the best general-purpose compression techniques. We have also simplified the implementation of SPIHT and reduced working storage requirements significantly compared to the original implementation. Our version of the algorithm allows the compression of large images such as mammograms with a standard PC. A research on the clinical applicability of the VQIC techniques in the context of very large digital mammogram archive is planned.

Acknowledgments

The authors would like to thank M.Sc J. Näppi for providing the microcalcification detection software.

References

- [1] C.N. Adams, A. Aiyer, B.J. Betts et al., "Image quality in lossy compressed digital mammograms", *Proc. 3rd Internat. Workshop in Digital Mammography*, Chicago, U.S.A, 1996.
- [2] V. Bhaskaran and K. Konstantinides, *Image and Video Compression Standards*, Kluwer Academic Publishers, Dordrecht, The Netherlands, 1995, Chapter 5.
- [3] J. Dengler, S. Behrens, J.F. Beraga, "Segmentation of microcalcifications in mammograms", *IEEE Trans. Medical Imaging*, Vol. 12, No. 4, December 1993.
- [4] B.J. Erickson, A. Manduca, K.R. Persons, "Clinical evaluation of wavelet compression of digitized chest X-rays", *Proc. SPIE 3031 Medical Imaging: Image Display*, Newport Beach, CA, 1997.
- [5] M. Giger, H. MacMahon, "Image processing and computer-aided diagnosis", *Radiologic Clinics of North America*, Vol. 34, No. 3, May 1996, pp. 565-596.
- [6] M. Hilton, B.D. Jawerth and A.N. Sengupta, "Compressing still and moving images", *Multimedia Systems*, Vol. 2, December 1994, pp. 218-227.
- [7] A. Manduca, A. Said, "Wavelet compression of medical images with set partitioning in hierarchical trees", *Proc. SPIE 2704 Medical Imaging: Image Display*, Newport Beach, CA, 1996.
- [8] P.J. Meer, R.L. Lagendijk, J. Biemond, "Local adaptive thresholding to reduce the bit rate in constant quality MPEG coding", *Proc. International picture coding symposium*, Melbourne, Australia, 1996
- [9] S.M. Perlmutter, P.C. Cosman, R.M. Gray et al., "Image quality in lossy compressed digital mammograms", *Signal Processing*, Vol. 59, No. 2, June 1997.
- [10] R. Plompen, J. Groenveld, F. Booman, D. Boekee, "An image knowledge based video codec for low bitrates", *Proc. SPIE 804 Advances in Image Processing*, 1987.
- [11] A. Said and W.A. Pearlman, "A new fast and efficient image codec based on set partitioning in hierarchical trees", *IEEE Trans. Circuits and Systems for Video Technology*, Vol. 6, June 1996, pp. 243-250.
- [12] J.M. Shapiro, "Embedded image coding using zerotrees of wavelet coefficients", *IEEE Trans. Signal Processing*, Vol. 31, No. 12, December 1993.
- [13] D. Shin, H. Wu, J. Liu, "A region of interest (ROI) based wavelet compression scheme for medical images", *Proc. SPIE 3031 Medical Imaging: Image Display*, Newport Beach, CA, 1997.
- [14] M. Vetterli and J. Kovačević: *Wavelets and Subband Coding*, Prentice Hall, Englewood Cliffs, NJ, 1995.

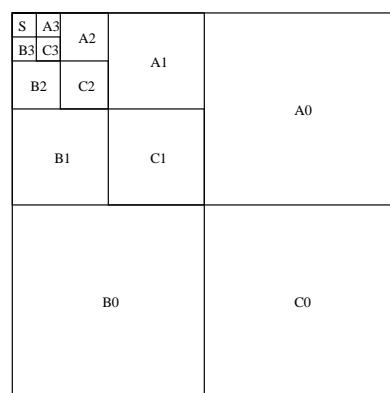


Fig. 1.

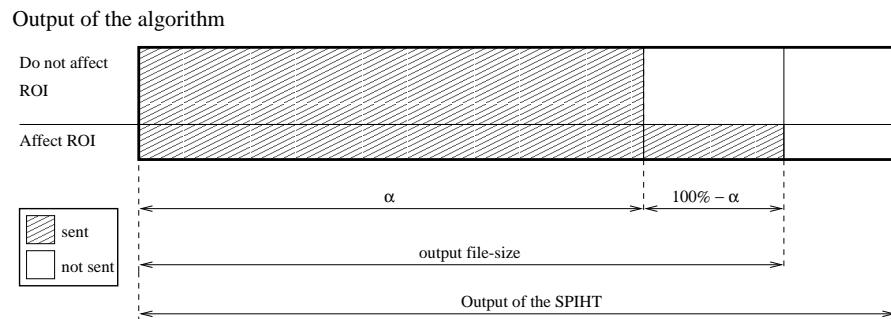
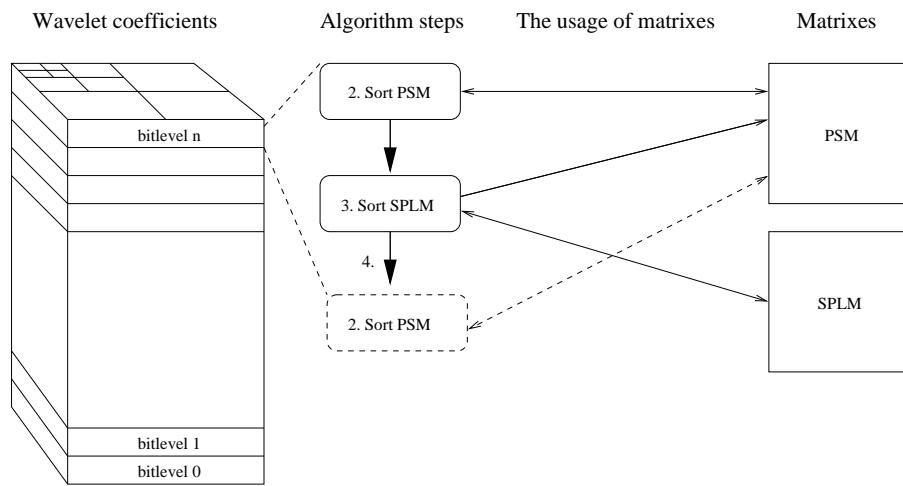


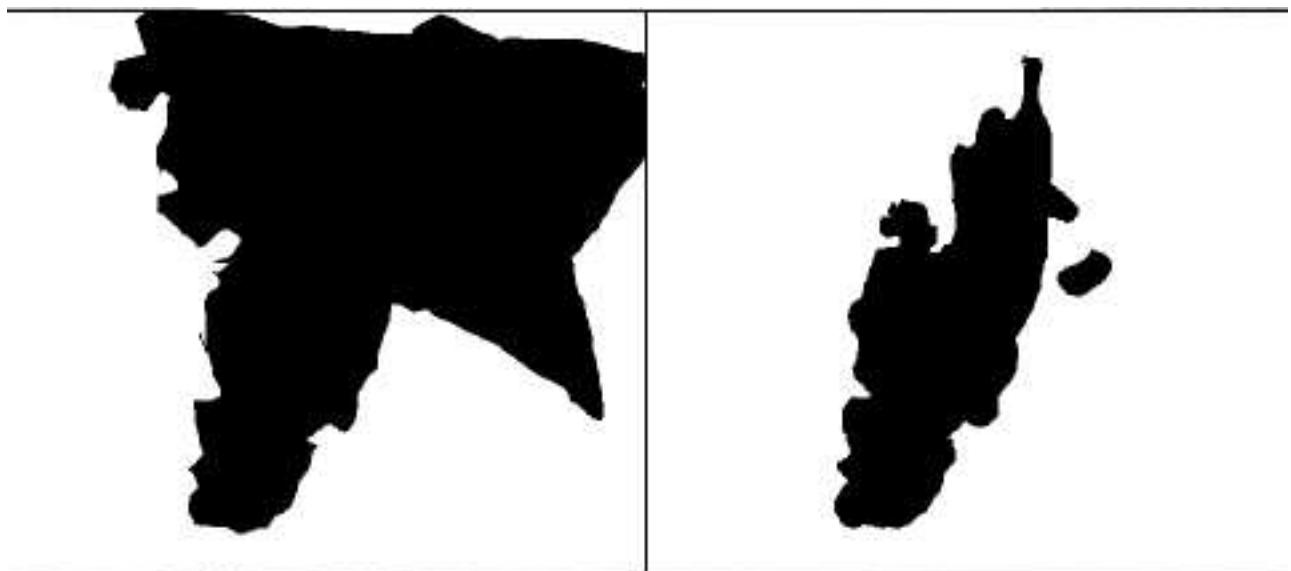
Fig. 2.



Fig. 3.



Fig. 4.



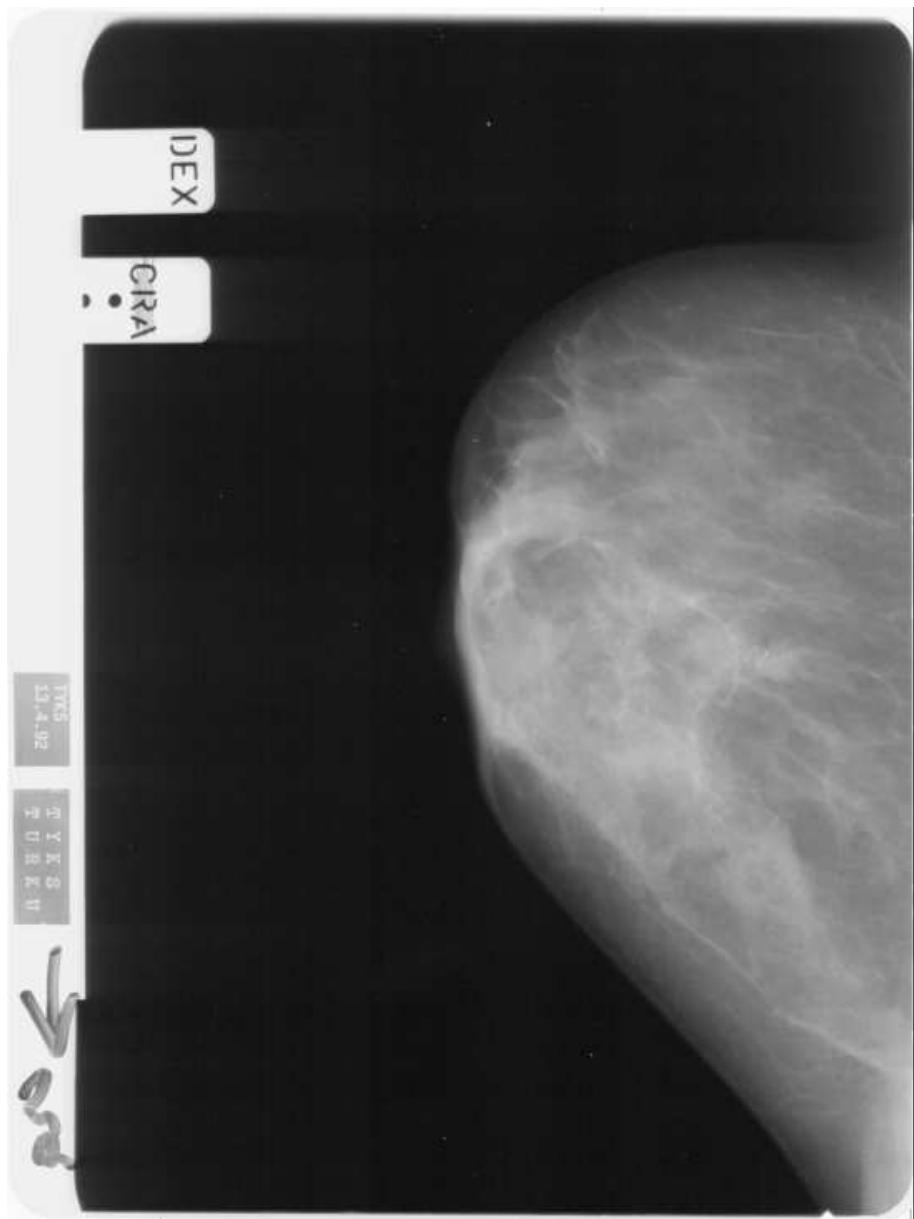


Fig. 6.

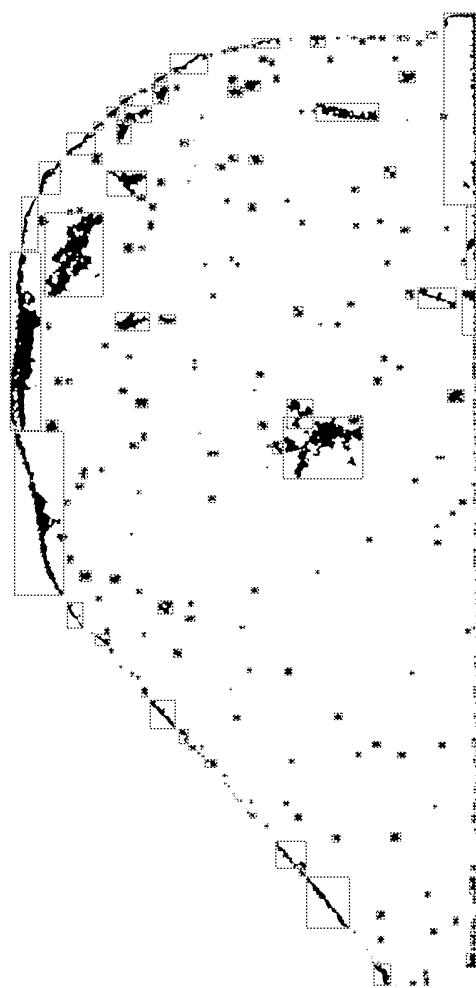


Fig. 7.

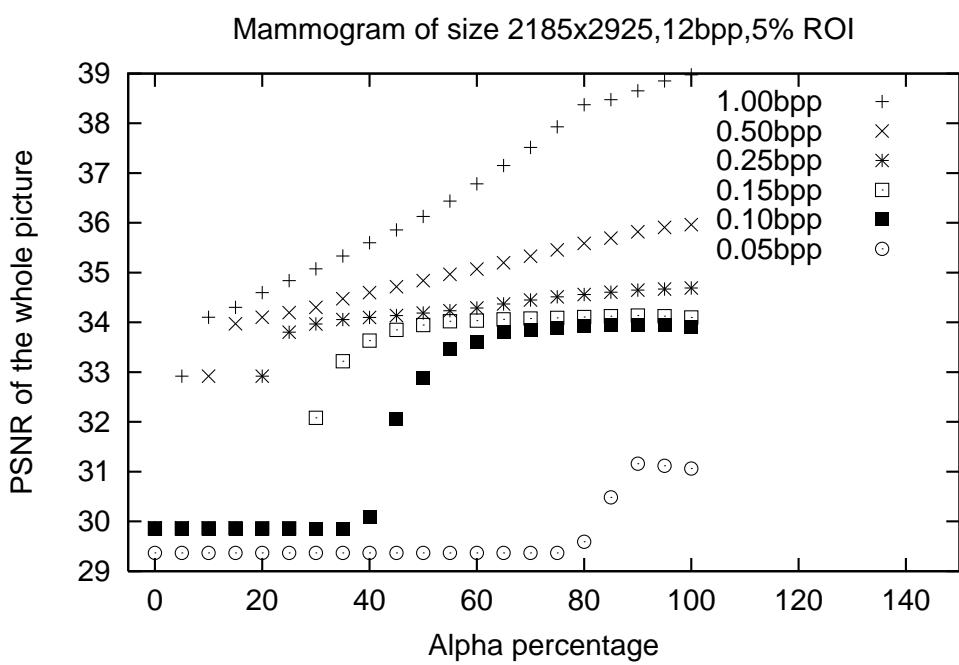


Fig. 8.

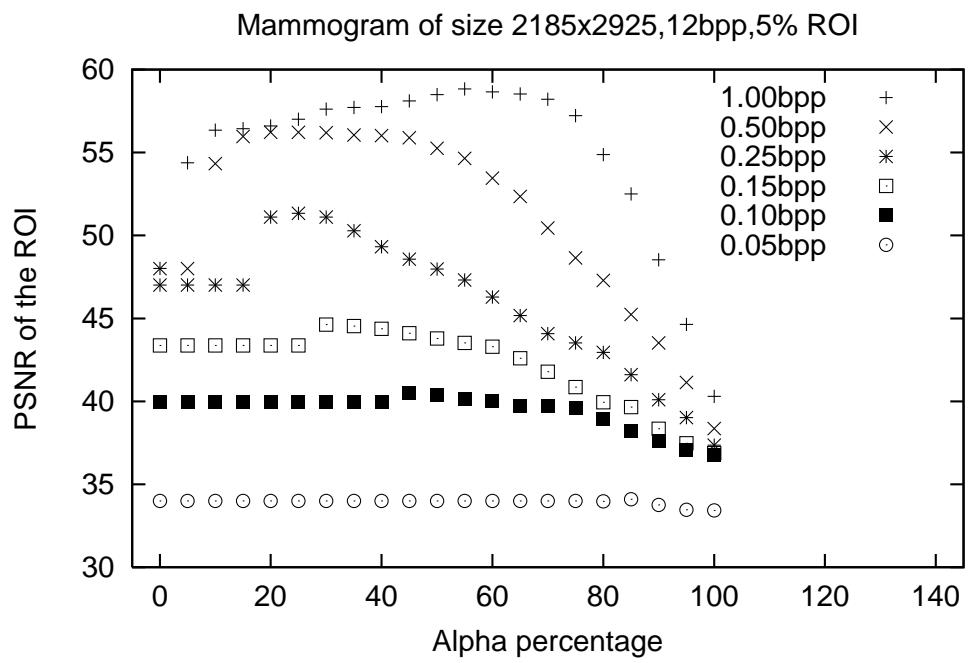


Fig. 9.

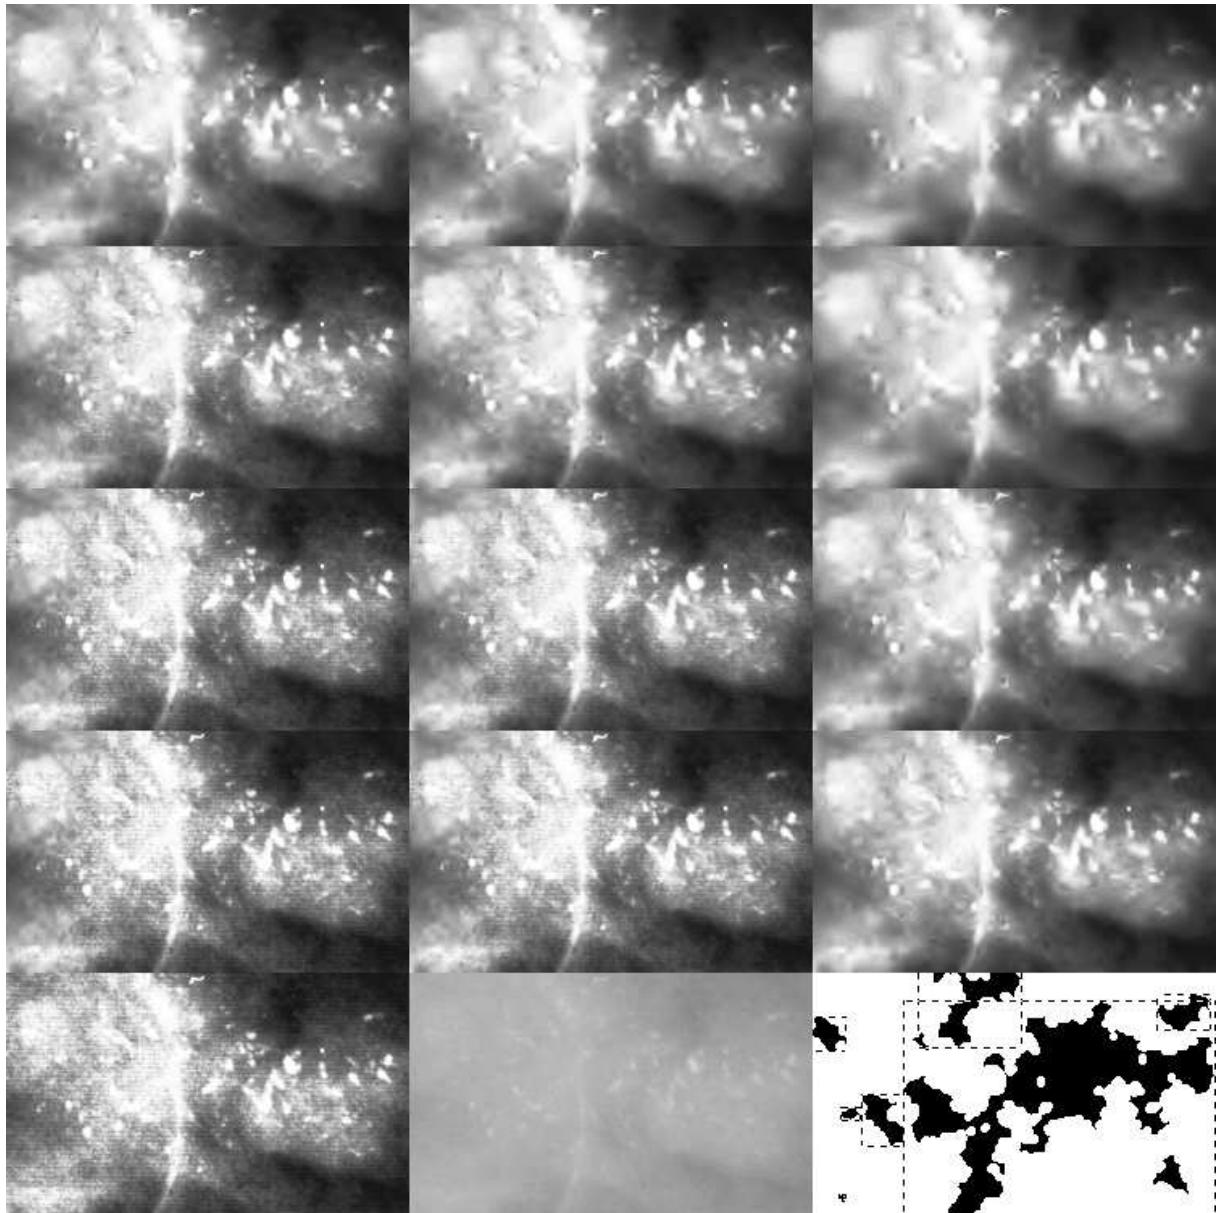


Fig. 10.

method	bpp	PSNR	PSNR in ROIs
GIF	4.456	∞	∞
JPEG	1.01	26.01	23.76
JPEG	0.50	23.09	20.21
JPEG	0.29	20.43	17.81
SPIHT	1.00	27.56	26.58
SPIHT	0.50	24.06	22.13
SPIHT	0.25	21.00	18.63
vqSPIHT $\alpha = 90\%$	1.00	27.26	39.85
vqSPIHT $\alpha = 90\%$	0.50	23.83	29.76
vqSPIHT $\alpha = 90\%$	0.25	20.75	23.41
vqSPIHT $\alpha = 80\%$	1.00	26.72	42.46
vqSPIHT $\alpha = 80\%$	0.50	23.40	34.36
vqSPIHT $\alpha = 80\%$	0.25	20.37	26.85
vqSPIHT $\alpha = 80\%$	0.15	18.55	22.13
vqSPIHT $\alpha = 80\%$	0.10	17.45	18.94
vqSPIHT $\alpha = 80\%$	0.05	15.15	14.55

Table 1

Comparison between SPIHT, vqSPIHT and JPEG for the comic image (Fig. 3).

Captions of figures:

Fig. 1. An example of a wavelet coefficient table, that contains three four-level pyramids: A , B and C . Scaling coefficients are located on the square noted by S .

Fig. 2. The vqSPIHT matrix implementation is illustrated on the upper half of the picture: The wavelet coefficients on the left are processed one bitlevel at a time. Each bitplane is processed in two steps (2. and 3.). The first step tests the significance of each coefficient that is labeled to be insignificant in PSM matrix and transmits the necessary information. The second step processes all the trees in SPLM and sets new points in PSM as *significant* or *insignificant*. The bottom part of the picture illustrates the thresholding of transmitted data. After the alpha limit has been reached, only bits of coefficients and decision information that affects the ROI is sent.

Fig. 3. The original comic test image with three ROIs marked.

Fig. 4. A region of the comic test image containing an ROI compressed with JPEG, SPIHT and vqSPIHT. See the table below for the explanation of sub-regions of the figure.

JPEG 1.00bpp	JPEG 0.50bpp	JPEG 0.25bpp
SPIHT 1.00bpp	SPIHT 0.50bpp	SPIHT 0.25bpp
vqSPIHT α 90% 1.00bpp	vqSPIHT α 90% 0.50bpp	vqSPIHT α 90% 0.25bpp
vqSPIHT α 80% 1.00bpp	vqSPIHT α 80% 0.50bpp	vqSPIHT α 80% 0.25bpp
vqSPIHT α 80% 0.15bpp	vqSPIHT α 80% 0.10bpp	ROI in original

Fig. 5. The two images on the left are ROI masks used in the compression of the two rightmost images, where ROI covers 16% (the upper image) and 46% of the image. The first grayscale image is compressed with without ROI ($\alpha = 100\%$), while 16% ROI ($\alpha = 50\%$) is used in the second image and 46% ROI ($\alpha = 50\%$) in the last image. All the images are compressed with same 0.25 bpp bitrate.

Fig. 6. Original mammogram test image.

Fig. 7. Micro-calcifications found in the test mammogram (shown as black), and the ROIs (dotted rectangles around micro-calcifications).

Fig. 8. PSNR of the whole mammogram for vqSPIHT as a function of α and bpp.

Fig. 9. PSNR of the ROIs in the mammogram for vqSPIHT as a function of α and bpp.

Fig. 10. A region of vqSPIHT compressed mammograms with different α and bpp rates. Bpp values on the columns from left to right are 1.00, 0.50 and 0.15. The values of α starting from the uppermost row are 100%, 90%, 70% and 50%. All the images have been histogram equalized to ease the evaluation. The three images in the last row from left to right are: the histogram equalized uncompressed region, original uncompressed region and a bit map of the detected micro-calcifications with the ROIs marked.

Table 1. Comparison between SPIHT, vqSPIHT and JPEG for the comic image (Fig. 3).

Footnotes

1. A common lossless image compression method.